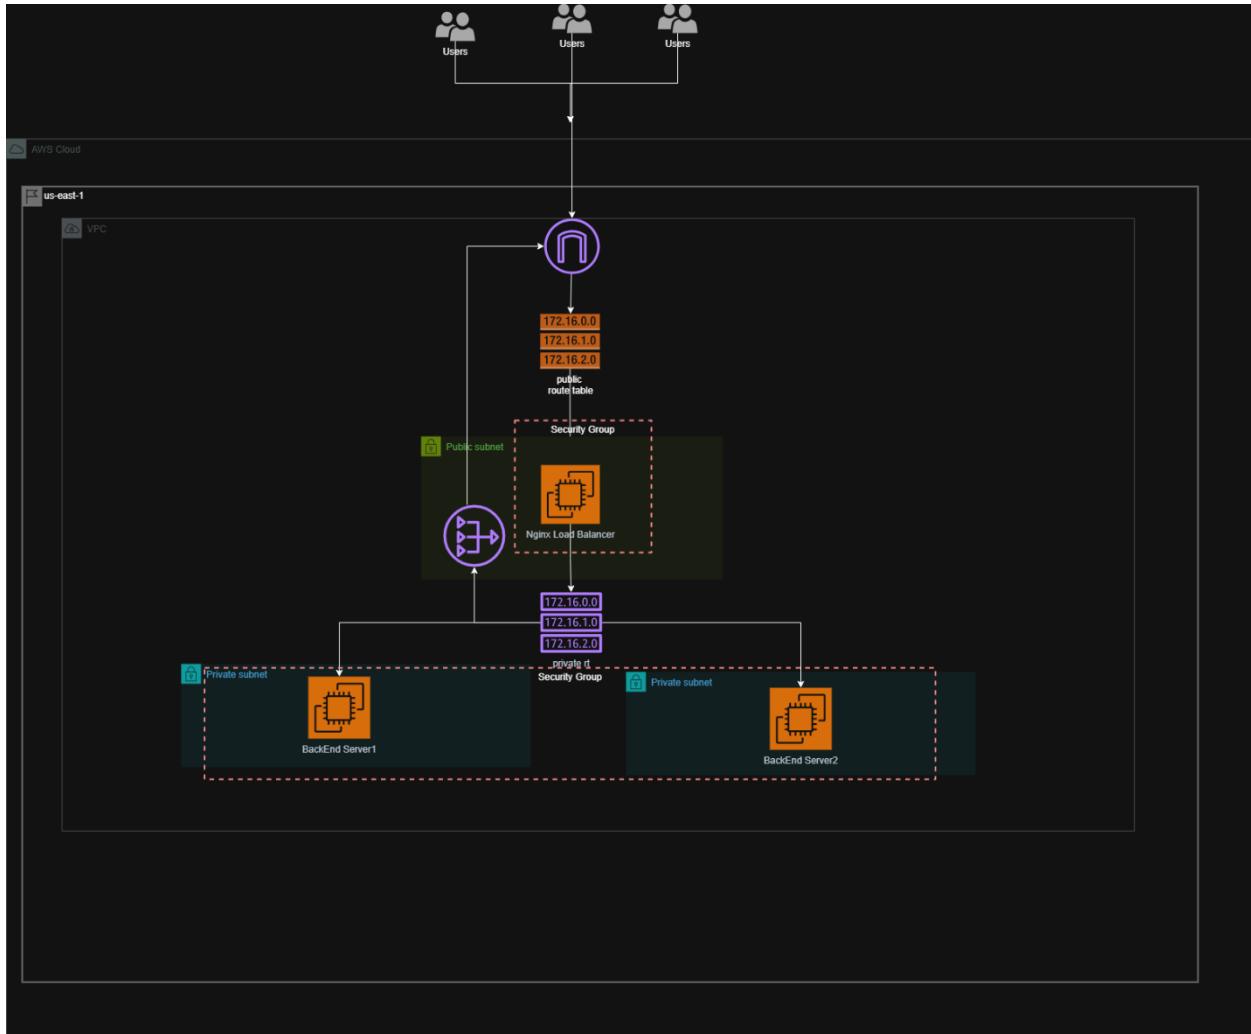

NGINX Load Balancer Setup on AWS EC2

Sardar Noor Ul Hassan (Cloudelligent)

1. Overview

This document explains how I deployed an NGINX-based load balancer on AWS to distribute traffic between two backend EC2 instances.

Both backend servers run in private subnets, and the load balancer runs in a public subnet. Traffic flows from the internet → NGINX LB → backend servers.



2. Prerequisites

- AWS account with permissions for EC2, VPC, and networking
- Basic understanding of Linux commands
- SSH key pair created and downloaded
- Custom VPC already configured with:
 - Public Subnet
 - Two Private Subnets
 - NAT Gateway
 - Internet Gateway
 - Security Groups for LB and backend

(VPC setup was already completed earlier.)

3. Backend EC2 Instances Setup

3.1 Launch Backend Server 1

- AMI: Ubuntu
- Instance Type: t3.micro
- Subnet: Private Subnet 1
- Auto-assign Public IP: Disabled
- Security Group: backend-servers-sg
- Key Pair: noor-key

3.2 Launch Backend Server 2

- AMI: Ubuntu
- Instance Type: t3.micro
- Subnet: Private Subnet 2
- Auto-assign Public IP: Disabled
- Security Group: backend-servers-sg
- Key Pair: noor-key

4. Install Web Server on Backend Instances

Since private EC2 instances cannot be accessed from the internet, I connected to them **through the Load Balancer EC2** using SSH.

Commands executed on each backend instance:

```
sudo apt update -y
```

```
sudo apt install apache2 -y
```

```
sudo systemctl enable apache2
```

```
sudo systemctl start apache2
```

```
echo "<h1>Hello from Backend Server1 (SardarNoor)</h1>" | sudo tee  
/var/www/html/index.html
```

For the second backend server, I replaced the message with:

```
Hello from Backend Server2 (SardarNoor)
```

Both backend servers were tested from the LB using:

```
curl http://  
Last login: Fri Nov 14 10:44:43 2023 from 119.75.100.68  
ubuntu@ip-10-0-1-202:~$ curl http://10.0.2.161  
curl http://10.0.3.204  
<h1>Hello from Backend Server1 (SardarNoor)</h1>  
<h1>Hello from Backend Server2 (SardarNoor)</h1>  
ubuntu@ip-10-0-1-202:~$
```

Both returned the expected HTML output.

5. Load Balancer EC2 Setup

5.1 Launch LB Instance

- AMI: Ubuntu
- Instance Type: t2.micro
- Subnet: Public Subnet
- Auto-assign Public IP: Enabled
- Security Group: lb-sg
 - HTTP allowed from anywhere (0.0.0.0/0)

- SSH allowed from My IP

5.2 Install NGINX on LB Instance

```
sudo apt update -y
```

```
sudo apt install nginx -y
```

```
sudo systemctl enable nginx
```

```
sudo systemctl start nginx
```

```
NGINX service confirmed running: sudo systemctl status nginx
```

6. Configure NGINX for Load Balancing

I edited the main NGINX configuration file:

```
sudo nano /etc/nginx/nginx.conf
```

Under the **http {}** block, I added:

```
http {  
    upstream backend_servers {  
        server 10.0.2.161;  
        server 10.0.3.204;  
    }  
}
```

```
server {  
    listen 80;  
  
    location / {  
        proxy_pass http://backend_servers;  
    }  
}
```

- ```
}
```
- 10.0.2.161 = Backend Server 1 (private IP)
  - 10.0.3.204 = Backend Server 2 (private IP)

Then tested and restarted NGINX:

```
sudo nginx -t
```

```
sudo systemctl restart nginx
```

```
http {
 upstream backend_servers {
 server 10.0.2.161;
 server 10.0.3.204;
 }

 server {
 listen 80;
 nginxlb http://35.165.79.43;

 location / {
 proxy_pass http://backend_servers;
 }
 }
}
```

---

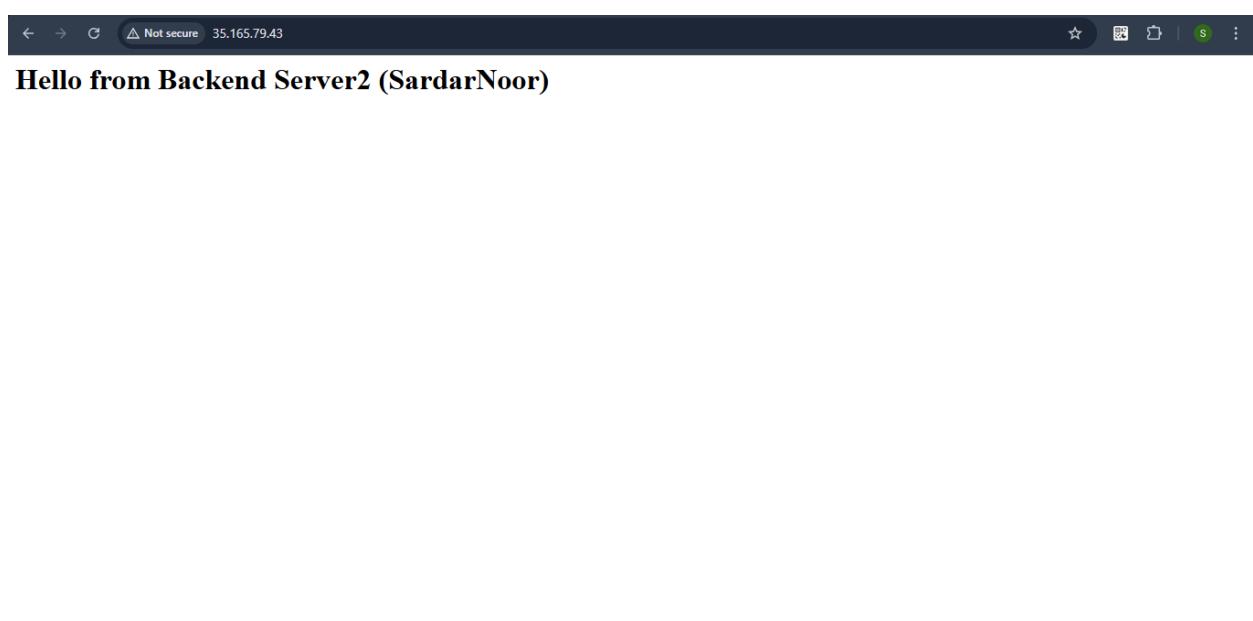
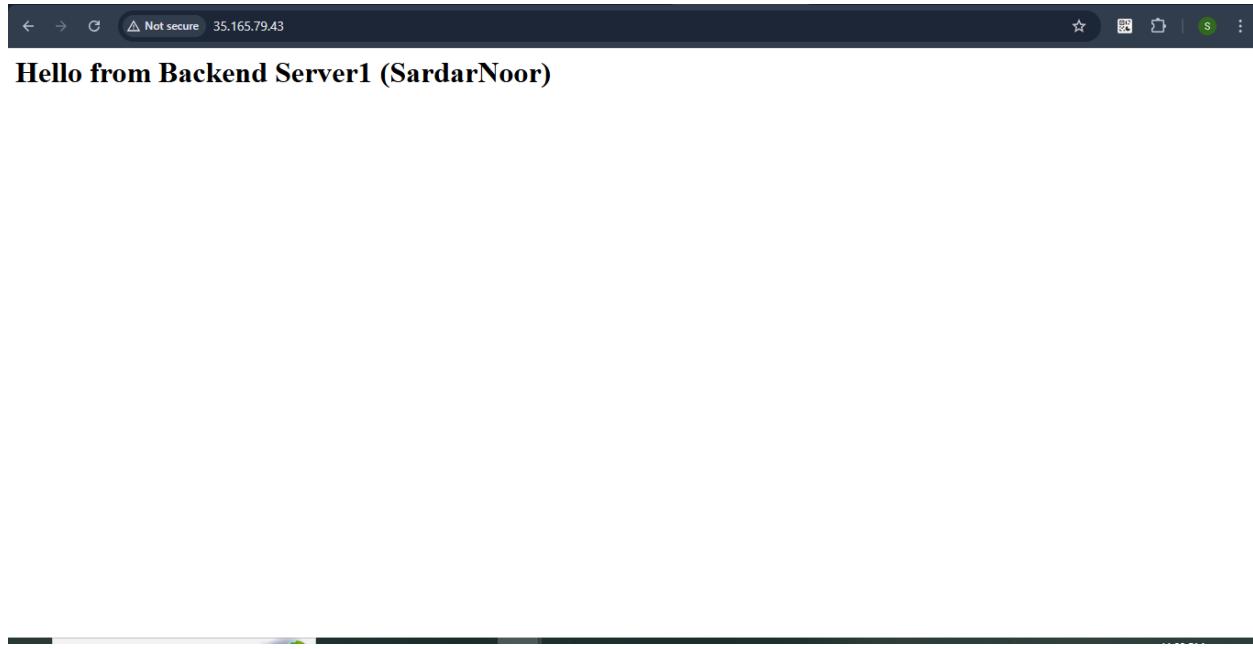
## 7. Testing the Load Balancer

I opened the LB's public IP in the browser:

```
http://35.165.79.43
```

Then refreshed several times.

Expected behavior:



- One refresh shows: “Hello from Backend Server1 (SardarNoor)”
- Next refresh shows: “Hello from Backend Server2 (SardarNoor)”
- And so on...

This confirmed **round-robin load balancing** is working correctly.

I also tested with:

```
curl http://35.165.79.43
```

```
ubuntu@ip-10-0-1-202:~$ curl http://35.165.79.43
<h1>Hello from Backend Server1 (SardarNoor)</h1>
ubuntu@ip-10-0-1-202:~$ curl http://35.165.79.43
<h1>Hello from Backend Server2 (SardarNoor)</h1>
ubuntu@ip-10-0-1-202:~$
```

Responses alternated between both backend servers.

---

## 8. Security Groups Summary

### 8.1 LB Security Group

Inbound:

- HTTP (80) → 0.0.0.0/0
- SSH (22) → My IP

Outbound:

- All traffic allowed

### 8.2 Backend Servers Security Group

Inbound:

- HTTP (80) → Source: lb-sg
- SSH (22) → My IP (optional; used only for LB → backend SSH testing)

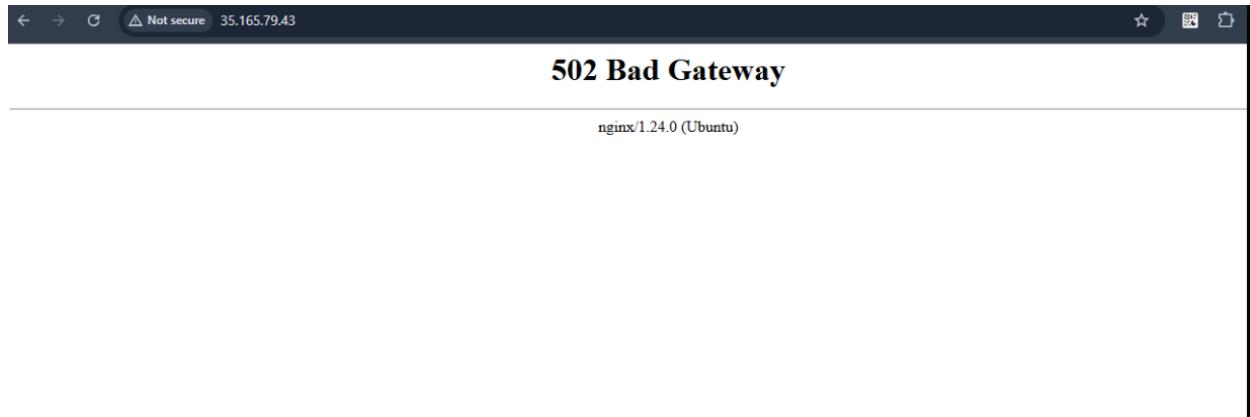
Outbound:

- All traffic allowed

## 11. Troubleshooting

During the setup, I encountered multiple issues while configuring the NGINX load balancer.

### 11.1 Issue: 502 Bad Gateway on Load Balancer



## Symptoms

- Browser showed 502 Bad Gateway
- curl: (7) Failed to connect

## Cause

NGINX was running, but the backend servers were **not responding on port 80**.

Apache was not installed/started on Backend Server 1, so NGINX had no valid upstream target.

## Solution

I connected to backend servers via the LB instance and manually installed & started Apache:

```
sudo apt update -y
```

```
sudo apt install apache2 -y
```

```
sudo systemctl enable apache2
```

```
sudo systemctl start apache2
```

```
echo "<h1>Hello from Backend ServerX (SardarNoor)</h1>" | sudo tee
/var/www/html/index.html
```

After installing Apache, the backend servers successfully responded:

```
curl http://10.0.2.161
```

```
curl http://10.0.3.204
```

502 error resolved.

## 11.2 Issue: Cannot SSH into Backend Servers

### Symptoms

- Direct SSH from my laptop failed (Connection timed out)
- LB → backend SSH failed with:
- Permission denied (publickey)

### Cause

Backend servers are in **private subnets** → cannot SSH directly.

Also, the LB instance did not have the key pair (noor-key.pem) required for backend SSH.

### Solution

1. Copied key to LB instance using SCP from laptop:

```
scp -i noor-key.pem noor-key.pem ubuntu@35.165.79.43:/home/ubuntu/
```

2. Fixed permissions on LB:

```
chmod 400 noor-key.pem
```

3. Added SSH rule in backend SG to allow SSH from LB-SG

4. Connected successfully:

```
ssh -i noor-key.pem ubuntu@10.0.2.161
```

```
ssh -i noor-key.pem ubuntu@10.0.3.204
```

SSH working → allowed backend debugging.

## 11.3 Issue: User Data Did Not Execute on Backend Servers

### Symptoms

- Apache not installed
- /var/www/html/index.html missing
- systemctl status apache2 returned:
- Unit apache2.service could not be found

### Cause

User-data script was missing the required shebang line:

```
#!/bin/bash
```

Also, user-data does NOT run on reboot ,it only runs on the first boot or after a full stop/start.

### Solution

I corrected the script and applied it again using:

```
#!/bin/bash
```

```
apt update -y
```

```
apt install apache2 -y
```

```
...
```

Then performed a **Stop → Start** (not reboot) to trigger user-data.

## 11.4 Issue: Load Balancer Not Switching Between Servers (Only Server 2 Response)

### Symptoms

- Browser always showed Backend Server 2
- No alternation between Server1 and Server2

### Cause

Backend Server 1 was not serving content due to missing Apache installation.

### Solution

After installing Apache on Backend Server 1 and verifying curl, NGINX started routing correctly:

Backend1 → Backend2 → Backend1 → Backend2

Round-robin balancing confirmed.

## 11.5 Issue: Initial Confusion with Route Tables

### Symptoms

At one point, I suspected route tables were incorrect.

### Cause

The backend servers were already using the correct private route table (NAT Gateway).  
This was not the actual issue ,the real issue was unpaid user-data.